

Saumya Bhati
202501040455

PRACTICE LAB ASSIGNMENT

4.1.1. PANDAS - SERIES CREATION AND MANIPULATION

Write a Python program that takes a list of numbers from the user, creates a Pandas Series from it, and then calculates the mean of even and odd numbers separately using the `groupby` and `mean()` operations.

Hint: You can group the Series based on whether each number is even or odd.

Input Format:

- The user should enter a list of numbers separated by space when prompted.

Output Format:

- The program should display the mean of even and odd numbers separately.
- Each mean value should be displayed with a label indicating whether it corresponds to even or odd numbers.

Refer to the sample test cases for better understanding regarding input and output format.

INPUT

```
import pandas as pd

# Take inputs from the user to create a list of numbers
numbers = list(map(int, input().split()))

# Create a Pandas series from the list of numbers
series=pd.Series(numbers)
# Grouping by even and odd numbers and calculating the mean
grouped =series.groupby(series%2==0).mean()

# Display the mean of even and odd numbers with labels
grouped.index = ['Even' if is_even else 'Odd' for is_even in
grouped.index]
print("Mean of even and odd numbers:")
print(grouped)
```

OUTPUT

Average time
0.029 s
29.00 ms

Maximum time
0.031 s
31.00 ms

✓ 3 out of 3 shown test case(s) passed
✓ 3 out of 3 hidden test case(s) passed

✓ Test case 1 31 ms Debug

Expected output
1 2 3 4 5 6 7 8 9 10
Mean of even and odd numbers:
Odd ···· 5.0
Even ···· 6.0
dtype: float64

Actual output
1 2 3 4 5 6 7 8 9 10
Mean of even and odd numbers:
Odd ···· 5.0
Even ···· 6.0
dtype: float64

✓ Test case 2 28 ms Debug

Expected output
4 4 4 6 6 2 2 2 10 12 14 16
Mean of even and odd numbers:
Even ···· 6.833333
dtype: float64

Actual output
4 4 4 6 6 2 2 2 10 12 14 16
Mean of even and odd numbers:
Even ···· 6.833333
dtype: float64

✓ Test case 3 28 ms Debug

Expected output
1 5 7 9 11 13 15 17 111 21 25
Mean of even and odd numbers:
Odd ···· 21.363636
dtype: float64

Actual output
1 5 7 9 11 13 15 17 111 21 25
Mean of even and odd numbers:
Odd ···· 21.363636
dtype: float64

4.1.2. DICTIONARY TO DATAFRAME

A dictionary of lists has been provided to you in the editor. Create a DataFrame from the dictionary of lists and perform the listed operations, then display the DataFrame before and after each manipulation.

Create the DataFrame:

- Convert the dictionary to a Pandas DataFrame.

Add a new row:

- Take inputs from the user for the new row data (name, age).
- Add the new row to the DataFrame.
- Display the DataFrame after adding the new row.

Modify a row:

- Modify a specific row by changing the age. Take the row index and new age value from the user.
- Display the DataFrame after modifying the row.

Delete a row:

- Take the row index to be deleted from the user.
- Remove the specified row.
- Display the DataFrame after deleting the row.

Add a new column:

- Add a column Gender with values taken from the user.
- Display the DataFrame after adding the new column.

Modify a column:

- Convert names to uppercase.
- Display the DataFrame after modifying the column.

Delete a column:

- Remove the Age column.
- Display the DataFrame after deleting the column.

INPUT

```
import pandas as pd

# Provided dictionary of lists
data = {
    'Name': ['Alice', 'Bob', 'Charlie'],
    'Age': [25, 30, 35],
}

# Convert the dictionary to a DataFrame
df = pd.DataFrame(data)

# Display the original DataFrame
print("Original DataFrame:")
print(df)
name=input("New name: ")
age=int(input("New age: "))
# Adding a new row
new_data={"Name":name,"Age":age}

df=df.append(new_data,ignore_index=True)
#df=pd.concat([df,pd.DataFrame([new_data]),ignore_index=True)

# Display the DataFrame after adding a new row
print("After adding a row:\n",df)

# Modifying a row
i=int(input("Index of row to modify: "))
new_age=int(input("New age: "))
df.loc[i,"Age"]=new_age

# Display the DataFrame after modifying a row
print("After modifying a row:")
print(df)

# Deleting a row
i=int(input("Index of row to delete: "))
df=df.drop(df.index[i]).reset_index(drop=True)
# Display the DataFrame after deleting a row
print("After deleting a row:")
print(df)

# Adding a new column
g_input=input("Enter genders separated by space: ")
genders=g_input.split()
df["Gender"]=genders

# Display the DataFrame after adding a new column
print("After adding a new column:")
print(df)

# Modifying a column
df["Name"]=df["Name"].str.upper()
print("After modifying a column:")
print(df)
```

OUTPUT

Average time
0.115 s
114.50 ms

Maximum time
0.119 s
119.00 ms

1 out of 1 shown test case(s) passed
1 out of 1 hidden test case(s) passed

Test case 1 119 ms Debug

Expected output

Actual output

Original DataFrame:
..... Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
New name: Susan
New age: 29
After adding a row:
..... Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to modify: 0
New age: 27
After modifying a row:
..... Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to delete: 3
After deleting a row:
..... Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
Enter genders separated by space: F M
After adding a new column:
..... Name Age Gender
0 Alice 27 F
1 Bob 30 M
2 Charlie 35 M
After modifying a column:
..... Name Age Gender
0 ALICE 27 F
1 BOB 30 M
2 CHARLIE 35 M
After deleting a column:
..... Name Gender
0 ALICE F
1 BOB M
2 CHARLIE M

Original DataFrame:
..... Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
New name: Susan
New age: 29
After adding a row:
..... Name Age
0 Alice 25
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to modify: 0
New age: 27
After modifying a row:
..... Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
3 Susan 29
Index of row to delete: 3
After deleting a row:
..... Name Age
0 Alice 27
1 Bob 30
2 Charlie 35
Enter genders separated by space: F M
After adding a new column:
..... Name Age Gender
0 Alice 27 F
1 Bob 30 M
2 Charlie 35 M
After modifying a column:
..... Name Age Gender
0 ALICE 27 F
1 BOB 30 M
2 CHARLIE 35 M
After deleting a column:
..... Name Gender
0 ALICE F
1 BOB M
2 CHARLIE M

4.1.3. STUDENT INFORMATION

Write a program to read a text file containing student information (name, age, and grade) using Pandas. Perform the following tasks:

- Display the first five rows of the dataframe.
- Calculate the average age of the students (limit the average age up to 2 decimal places).
- Filter out the students who have a grade above a certain threshold (consider the threshold grade is 'B').

INPUT

```
import pandas as pd

# Read the text file into a DataFrame
file = input()
data = pd.read_csv(file, sep="\s+", header=None, names=
["Name", "Age", "Grade"])

# write your code here..
df=pd.DataFrame(data)
print("First five rows:")
print(df.head(5))
print("Average age:",df["Age"].mean().round(2))
print("Students with a grade up to B")
filtered_df=df[df['Grade'].isin(['A', 'B'])]
print(filtered_df)
```

OUTPUT

Average time

0.034 s

34.00 ms

Maximum time

0.040 s

40.00 ms

1 out of 1 shown test case(s) passed

1 out of 1 hidden test case(s) passed

Test case 1

40 ms

Debug

Expected output

Actual output

studentdata.txt

studentdata.txt

First five rows:

First five rows:

Name

Age

Grade

Name

Age

Grade

0 John 25 A

0 John 25 A

1 Alin 22 B

1 Alin 22 B

2 Emma 24 A

2 Emma 24 A

3 Mike 23 C

3 Mike 23 C

4 Sony 21 B

4 Sony 21 B

Average age: 23.29

Average age: 23.29

Students with a grade up to B

Students with a grade up to B

Name

Age

Grade

Name

Age

Grade

0 John 25 A

0 John 25 A

1 Alin 22 B

1 Alin 22 B

2 Emma 24 A

2 Emma 24 A

4 Sony 21 B

4 Sony 21 B

5 Amia 26 A

5 Amia 26 A

LAB ASSIGNMENT

4.2.1. MONTH WITH THE HIGHEST TOTAL SALES

Write a Python program that takes the filename of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by Month and calculate the total sales for each month.
- Find the month with the highest total sales and display it.
- Also, display the total sales for the best month.

Sample Data:

Date,Product ,Quantity,Price,City

2025-01-01,Product A,5,20,New York

2025-01-01,Product B,3,15,Los Angeles

2025-01-02,Product A,7,20,New York

2025-01-02,Product C,4,30,Chicago

2025-01-03,Product B,2,15,Chicago

2025-01-03,Product A,8,20,Los Angeles

2025-01-04,Product C,6,30,New York

2025-01-04,Product B,5,15,Los Angeles

2025-01-05,Product A,3,20,Chicago

2025-01-05,Product C,10,30,Los Angeles

INPUT

```
import pandas as pd

# Prompt the user for the file name
file_name = input()

# Load the data
df = pd.read_csv(file_name)

df['Date'] = pd.to_datetime(df['Date'])

# Calculate total sales for each transaction
df['Total_Sales'] = df['Quantity'] * df['Price']

# Extract month and year from Date
df['Month'] = df['Date'].dt.to_period('M') # Creates YYYY-MM
format

# Find the month with the highest total sales
monthly_sales = df.groupby('Month')['Total_Sales'].sum()
best_month = monthly_sales.idxmax()
highest_sales = monthly_sales.max()

print(f"Best month: {best_month}")
print(f"Total sales: ${highest_sales:.2f}")
```

OUTPUT

Average time
0.013 s
13.00 ms

Maximum time
0.029 s
29.00 ms

✓ 1 out of 1 shown test case(s) passed
✓ 2 out of 2 hidden test case(s) passed

✓ Test case 1 29 ms Debug

Expected output

sales_data.csv

Best month: 2025-01↵

Total sales: \$1210.00↵

Actual output

sales_data.csv

Best month: 2025-01↵

Total sales: \$1210.00↵

4.2.2. BEST SELLING PRODUCT

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Find the product that sold the most in terms of quantity sold.
- Display the product that sold the most and the total quantity sold for that product.

Sample Data:

Date,Product ,Quantity,Price,City
2025-01-01,Product A,5,20,New York
2025-01-01,Product B,3,15,Los Angeles
2025-01-02,Product A,7,20,New York
2025-01-02,Product C,4,30,Chicago
2025-01-03,Product B,2,15,Chicago
2025-01-03,Product A,8,20,Los Angeles
2025-01-04,Product C,6,30,New York
2025-01-04,Product B,5,15,Los Angeles
2025-01-05,Product A,3,20,Chicago
2025-01-05,Product C,10,30,Los Angeles

INPUT

```
import pandas as pd

# Prompt the user for the file name
file_name = input()

# Load the data
df = pd.read_csv(file_name)

product_quantity = df.groupby('Product')['Quantity'].sum()

best_product = product_quantity.idxmax()
highest_quantity = product_quantity.max()

# Display the result
print(f"Best selling product: {best_product}")
print(f"Total quantity sold: {highest_quantity}")
```

OUTPUT

Average time
0.010 s
10.33 ms

Maximum time
0.025 s
25.00 ms

✓ 1 out of 1 shown test case(s) passed
✓ 2 out of 2 hidden test case(s) passed

✓ Test case 1 25 ms Debug

Expected output

sales_data.csv

Best selling product: Product A↵

Total quantity sold: 23↵

Actual output

sales_data.csv

Best selling product: Product A↵

Total quantity sold: 23↵

4.2.3. CITY THAT SOLD THE MOST PRODUCTS

Write a Python program that takes the file name of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the columns: Date, Product, Quantity, Price, and City.
- Group the data by City and calculate the total quantity of products sold for each city.
- Find the city that sold the most products (based on the total quantity sold).

Sample Data:

Date,Product ,Quantity,Price,City
2025-01-01,Product A,5,20,New York
2025-01-01,Product B,3,15,Los Angeles
2025-01-02,Product A,7,20,New York
2025-01-02,Product C,4,30,Chicago
2025-01-03,Product B,2,15,Chicago
2025-01-03,Product A,8,20,Los Angeles
2025-01-04,Product C,6,30,New York
2025-01-04,Product B,5,15,Los Angeles
2025-01-05,Product A,3,20,Chicago
2025-01-05,Product C,10,30,Los Angeles

INPUT

```
import pandas as pd

# Prompt the user for the file name
file_name = input()

# Load the data
df = pd.read_csv(file_name)

city_sales = df.groupby('City')['Quantity'].sum()

best_city = city_sales.idxmax()

# Display the result
print(f"City sold the most products: {best_city}")
```

OUTPUT

Average time
0.010 s
10.33 ms

Maximum time
0.025 s
25.00 ms

✓ 1 out of 1 shown test case(s) passed
✓ 2 out of 2 hidden test case(s) passed

✓ Test case 1 25 ms Debug

Expected output
sales_data.csv
City sold the most products: Los Angeles

Actual output
sales_data.csv
City sold the most products: Los Angeles

4.2.4. MOST FREQUENTLY SOLD PRODUCT PAIRS

Write a Python program that takes the filename of a CSV file as input, reads the data, and performs the following operations:

- The CSV file contains the following columns: Date, Product, Quantity, Price, and City.
- For each date, find all pairs of products that were sold together (i.e., two products sold on the same date).
- Output the product pair/s that was sold most frequently.

Sample Data:

Date,Product ,Quantity,Price,City

2025-01-01,Product A,5,20,New York

2025-01-01,Product B,3,15,Los Angeles

2025-01-02,Product A,7,20,New York

2025-01-02,Product C,4,30,Chicago

2025-01-03,Product B,2,15,Chicago

2025-01-03,Product A,8,20,Los Angeles

2025-01-04,Product C,6,30,New York

2025-01-04,Product B,5,15,Los Angeles

2025-01-05,Product A,3,20,Chicago

2025-01-05,Product C,10,30,Los Angeles

Explanation:

Transactions:

- 2025-01-01: Product A, Product B
- 2025-01-02: Product A, Product C
- 2025-01-03: Product B, Product A
- 2025-01-04: Product C, Product B
- 2025-01-05: Product A, Product C
-

Now, let's count how often the pairs of products appear together:

- Product A and Product B: Appear in transactions on 2025-01-01 and 2025-01-03.
- Product A and Product C: Appear in transactions on 2025-01-02 and 2025-01-05.
- Product B and Product C: Appears in transactions on 2025-01-04.

Most Frequent Product Combinations:

- Product A and Product B (2 times)
- Product A and Product C (2 times)

INPUT

```
import pandas as pd
from itertools import combinations
from collections import Counter

# Prompt user to input the file name
file_name = input()

# Read data from the specified CSV file
df = pd.read_csv(file_name)

daily_transactions = df.groupby('Date')['Product'].apply(list)
# Initialize a Counter to store the frequency of each pair
pair_counts = Counter()
# 2. Iterate through each date's transaction
for products in daily_transactions:
# Sort the products to ensure (A, B) is treated the same as
(B, A)
    products = sorted(products)

    pairs = list(combinations(products, 2))
# Update the counter with these pairs
    pair_counts.update(pairs)

if pair_counts:
    max_count = max(pair_counts.values())
    for pair, count in pair_counts.items():
        if count == max_count:
            print(f"{pair[0]} and {pair[1]}: {count} times")
else:
    print("No product pairs found.")
```

OUTPUT

Average time

0.010 s

9.67 ms



Maximum time

0.022 s

22.00 ms



✓ 1 out of 1 shown test case(s) passed

✓ 2 out of 2 hidden test case(s) passed

✓ Test case 1 22 ms

Debug





^

Expected output

sales_data.csv

Product A and Product B: 2 times

Product A and Product C: 2 times

Actual output

sales_data.csv

Product A and Product B: 2 times

Product A and Product C: 2 times

4.2.5. TITANIC DATASET ANALYSIS AND DATA CLEANING

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset. For each question, perform necessary data cleaning, transformations, and calculations as required.

- 1. Display the first 5 rows of the dataset.
- 2. Display the last 5 rows of the dataset.
- 3. Get the shape of the dataset (number of rows and columns).
- 4. Get a summary of the dataset (using .info()).
- 5. Get basic statistics (mean, standard deviation, etc.) of the dataset using .describe().
- 6. Check for missing values and display the count of missing values for each column.
- 7. Fill missing values in the 'Age' column with the median age.
- 8. Fill missing values in the 'Embarked' column with the most frequent value (mode()).
- 9. Drop the 'Cabin' column due to many missing values.
- 10. Create a new column, 'FamilySize' by adding the 'SibSp' and 'Parch' columns.

The Titanic dataset contains columns as shown below,

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:
PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

INPUT

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')

# 1. Display the first 5 rows of the dataset
print(data.head())
# 2. Display the last 5 rows of the dataset
print(data.tail())
# 3. Get the shape of the dataset
print(data.shape)
# 4. Get a summary of the dataset (info)
print(data.info())
# 5. Get basic statistics of the dataset
print(data.describe())
# 6. Check for missing values
print(data.isnull().sum())
# 7. Fill missing values in the 'Age' column with the median age
data['Age'].fillna(data['Age'].median(), inplace = True)
# 8. Fill missing values in the 'Embarked' column with the mode
data['Embarked'].fillna(data['Embarked'].mode()[0],inplace=True)
# 9. Drop the 'Cabin' column due to many missing values
data.drop(columns='Cabin',inplace = True)
# 10. Create a new column 'FamilySize' by adding 'SibSp' and 'Parch'
data["FamilySize"] =data['SibSp']+data['Parch']
```

OUTPUT

Average time
0.069 s
69.00 ms

Maximum time
0.069 s
69.00 ms

1 out of 1 shown test case(s) passed

Test case 1 69 ms Debug

Expected output

PassengerId	Survived	Pclass	Fare	Cabin	Embarked
0	1	0	7.2500	NaN	S
1	2	1	71.2833	C85	C
2	3	1	7.9250	NaN	S
3	4	1	53.1000	C123	S
4	5	0	8.0500	NaN	S

[5 rows x 12 columns]

PassengerId	Survived	Pclass	Fare	Cabin	Embarked
886	887	0	13.00	NaN	S
887	888	1	30.00	B42	S
888	889	0	23.45	NaN	S
889	890	1	30.00	C148	C
890	891	0	7.75	NaN	Q

[5 rows x 12 columns]

(891, 12)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 891 entries, 0 to 890

Data columns (total 12 columns):

Column	Non-Null Count	Dtype
0 PassengerId	891 non-null	int64

[1 Survived 891 non-null int64

[2 Pclass 891 non-null int64

[3 Name 891 non-null object

[4 Sex 891 non-null object

[5 Age 714 non-null float64

[6 SibSp 891 non-null int64

[7 Parch 891 non-null int64

[8 Ticket 891 non-null object

[9 Fare 891 non-null float64

[10 Cabin 204 non-null object

[11 Embarked 889 non-null object

dtypes: float64(2), int64(5), object(5)

memory usage: 66.2+ KB

None

PassengerId	Survived	Pclass	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	0.523008	0.381594
std	257.353842	0.486592	0.836071	1.102743	0.806057
min	1.000000	0.000000	1.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	1.000000	31.000000
max	891.000000	1.000000	3.000000	8.000000	512.329200

[8 rows x 7 columns]

PassengerId	Survived	Pclass	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000

Actual output

PassengerId	Survived	Pclass	Fare	Cabin	Embarked
0	1	0	7.2500	NaN	S
1	2	1	71.2833	C85	C
2	3	1	7.9250	NaN	S
3	4	1	53.1000	C123	S
4	5	0	8.0500	NaN	S

[5 rows x 12 columns]

PassengerId	Survived	Pclass	Fare	Cabin	Embarked
886	887	0	13.00	NaN	S
887	888	1	30.00	B42	S
888	889	0	23.45	NaN	S
889	890	1	30.00	C148	C
890	891	0	7.75	NaN	Q

[5 rows x 12 columns]

(891, 12)

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 891 entries, 0 to 890

Data columns (total 12 columns):

Column	Non-Null Count	Dtype
0 PassengerId	891 non-null	int64
1 Survived	891 non-null	int64
2 Pclass	891 non-null	int64
3 Name	891 non-null	object
4 Sex	891 non-null	object
5 Age	714 non-null	float64
6 SibSp	891 non-null	int64
7 Parch	891 non-null	int64
8 Ticket	891 non-null	object
9 Fare	891 non-null	float64
10 Cabin	204 non-null	object
11 Embarked	889 non-null	object

dtypes: float64(2), int64(5), object(5)

memory usage: 66.2+ KB

None

PassengerId	Survived	Pclass	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	891.000000	891.000000

4.2.6. TITANIC DATASET ANALYSIS AND DATA CLEANING - 2

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

- 1. Create a new column 'IsAlone' which is 1 if the passenger is alone (FamilySize=0), otherwise 0.
- 2. Convert the 'Sex' column to numeric values (male: 0, female: 1).
- 3. One-hot encode the 'Embarked' column, dropping the first category.
- 4. Get the mean age of passengers.
- 5. Get the median fare of passengers.
- 6. Get the number of passengers by class.
- 7. Get the number of passengers by gender.
- 8. Get the number of passengers by survival status.
- 9. Calculate the survival rate of passengers.
- 10. Calculate the survival rate by gender.

The Titanic dataset contains columns as shown below,

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:
PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

INPUT

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']

# Create a new column 'IsAlone' which is 1 if the passenger
is alone

data['IsAlone'] = np.where(data['FamilySize'] == 0, 1, 0)

# 2. Convert 'Sex' to numeric (male: 0, female: 1)
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
# 3. One-hot encode the 'Embarked' column
data = pd.get_dummies(data, columns=['Embarked'])
# 4. Get the mean age of passengers
mean_age = data['Age'].mean()
print(mean_age)
# 5. Get the median fare of passengers
median_fare = data['Fare'].median()
print(median_fare)
# 6. Get the number of passengers by class
print(data['Pclass'].value_counts())
# 7. Get the number of passengers by gender
print(data['Sex'].value_counts())
# 8. Get the number of passengers by survival status
print(data['Survived'].value_counts())
# 9. Calculate the overall survival rate
survival_rate = data['Survived'].mean()
print(format(survival_rate))
# 10. Calculate the survival rate by gender
survival_by_gender = data.groupby('Sex')['Survived'].mean()
print(survival_by_gender)
```

OUTPUT

Average time
0.032 s
32.00 ms

Maximum time
0.032 s
32.00 ms

1 out of 1 shown test case(s) passed

Test case 1 32 ms Debug

Expected output

29.69911764705882

14.4542

3...491

1...216

2...184

Name: Pclass, dtype: int64

0...577

1...314

Name: Sex, dtype: int64

0...549

1...342

Name: Survived, dtype: int64

0.3838383838383838

Sex

0...0.188908

1...0.742038

Name: Survived, dtype: float64

Actual output

29.69911764705882

14.4542

3...491

1...216

2...184

2...184

Name: Pclass, dtype: int64

0...577

1...314

Name: Sex, dtype: int64

0...549

1...342

Name: Survived, dtype: int64

0.3838383838383838

Sex

0...0.188908

1...0.742038

Name: Survived, dtype: float64

4.2.7. TITANIC DATASET ANALYSIS AND DATA CLEANING - 3

You are provided with the titanic dataset containing information about passengers on the Titanic. Your task is to write python code to answer the following questions based on the dataset.

- 1. Calculate the survival rate by class.
- 2. Calculate the survival rate by embarkation location (Embarked_S).
- 3. Calculate the survival rate by family size (FamilySize).
- 4. Calculate the survival rate by being alone (IsAlone).
- 5. Get the average fare by passenger class (Pclass).
- 6. Get the average age by passenger class (Pclass).
- 7. Get the average age by survival status (Survived).
- 8. Get the average fare by survival status (Survived).
- 9. Get the number of survivors by class (Pclass).
- 10. Get the number of non-survivors by class (Pclass).

The Titanic dataset contains columns as shown below,

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:
PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

INPUT

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data['FamilySize'] = data['SibSp'] + data['Parch']
data['IsAlone'] = np.where(data['FamilySize'] > 0, 0, 1)
data = pd.get_dummies(data, columns=['Embarked'],
drop_first=True)

# 1. Calculate the survival rate by class
print(data.groupby('Pclass')['Survived'].mean())
# 2. Calculate the survival rate by embarkation location
(Embarked_S)
print(data.groupby('Embarked_S')['Survived'].mean())
# 3. Calculate the survival rate by family size
print(data.groupby('FamilySize')['Survived'].mean())
# 4. Calculate the survival rate by being alone
print(data.groupby('IsAlone')['Survived'].mean())
# 5. Get the average fare by passenger class
print(data.groupby('Pclass')['Fare'].mean())
# 6. Get the average age by passenger class
print(data.groupby('Pclass')['Age'].mean())
# 7. Get the average age by survival status
print(data.groupby('Survived')['Age'].mean())
# 8. Get the average fare by survival status
print(data.groupby('Survived')['Fare'].mean())
# 9. Get the number of survivors by class
# Filter for Survived=1, then count occurrences of Pclass
print(data[data['Survived'] == 1]['Pclass'].value_counts())
# 10. Get the number of non-survivors by class
# Filter for Survived=0, then count occurrences of Pclass
print(data[data['Survived'] == 0]['Pclass'].value_counts())
```

OUTPUT

Average time
0.038 s
38.00 ms

Maximum time
0.038 s
38.00 ms

1 out of 1 shown test case(s) passed

Test case 1 38 ms Debug

Expected output

Pclass

1...0.629630

2...0.472826

3...0.242363

Name: Survived, dtype: float64

Embarked_S

0...0.506073

1...0.336957

Name: Survived, dtype: float64

FamilySize

0...0.303538

1...0.552795

2...0.578431

3...0.724138

4...0.200000

5...0.136364

6...0.333333

7...0.000000

10...0.000000

Name: Survived, dtype: float64

IsAlone

0...0.505650

1...0.303538

Name: Survived, dtype: float64

Pclass

1...84.154687

2...20.662183

3...13.675550

Name: Fare, dtype: float64

Age

1...38.233441

2...29.877630

3...25.140620

Name: Age, dtype: float64

Survived

0...30.626179

1...28.343690

Name: Age, dtype: float64

Survived

0...22.117887

1...48.395408

Name: Fare, dtype: float64

Survived

0...136

3...119

2...87

Name: Pclass, dtype: int64

3...372

2...97

1...80

Name: Pclass, dtype: int64

Actual output

Pclass

1...0.629630

2...0.472826

3...0.242363

Name: Survived, dtype: float64

Embarked_S

0...0.506073

1...0.336957

Name: Survived, dtype: float64

FamilySize

0...0.303538

1...0.552795

2...0.578431

3...0.724138

4...0.200000

5...0.136364

6...0.333333

7...0.000000

10...0.000000

Name: Survived, dtype: float64

IsAlone

0...0.505650

1...0.303538

Name: Survived, dtype: float64

Pclass

1...84.154687

2...20.662183

3...13.675550

Name: Fare, dtype: float64

Age

1...38.233441

2...29.877630

3...25.140620

Name: Age, dtype: float64

Survived

0...30.626179

1...28.343690

Name: Age, dtype: float64

Survived

0...22.117887

1...48.395408

Name: Fare, dtype: float64

Survived

0...136

3...119

2...87

Name: Pclass, dtype: int64

3...372

2...97

1...80

4.2.8. TITANIC DATASET ANALYSIS AND DATA CLEANING - 4

You are provided with the Titanic dataset containing information about passengers on the Titanic. Your task is to write Python code to answer the following questions based on the dataset.

- 1. Get the number of survivors by gender (Sex).
- 2. Get the number of non-survivors by gender (Sex).
- 3. Get the number of survivors by embarkation location (Embarked_S).
- 4. Get the number of non-survivors by embarkation location (Embarked_S).
- 5. Calculate the percentage of children (Age < 18) who survived.
- 6. Calculate the percentage of adults (Age >= 18) who survived.
- 7. Get the median age of survivors.
- 8. Get the median age of non-survivors.
- 9. Get the median fare of survivors.
- 10. Get the median fare of non-survivors.

The Titanic dataset contains columns as shown below,

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked

Sample Data:
PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, Embarked
1,0,3,"Braund, Mr. Owen Harris",male,22,1,0,A/5 21171,7.25,,S
2,1,1,"Cumings, Mrs. John Bradley (Florence Briggs Thayer)",female,38,1,0,PC 17599,71.2833,C85,C
3,1,3,"Heikkinen, Miss. Laina",female,26,0,0,STON/O2. 3101282,7.925,,S
4,1,1,"Futrelle, Mrs. Jacques Heath (Lily May Peel)",female,35,1,0,113803,53.1,C123,S
5,0,3,"Allen, Mr. William Henry",male,35,0,0,373450,8.05,,S
6,0,3,"Moran, Mr. James",male,,0,0,330877,8.4583,,Q
7,0,1,"McCarthy, Mr. Timothy J",male,54,0,0,17463,51.8625,E46,S
8,0,3,"Palsson, Master. Gosta Leonard",male,2,3,1,349909,21.075,,S
9,1,3,"Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)",female,27,0,2,347742,11.1333,,S
10,1,2,"Nasser, Mrs. Nicholas (Adele Achem)",female,14,1,0,237736,30.0708,,C

INPUT

```
import pandas as pd
import numpy as np

# Load the Titanic dataset
data = pd.read_csv('Titanic-Dataset.csv')
data = pd.get_dummies(data, columns=['Embarked'],
drop_first=True)

# 1. Get the number of survivors by gender
print(data[data['Survived']==1]['Sex'].value_counts())
# 2. Get the number of non-survivors by gender
print(data[data['Survived']==0]['Sex'].value_counts())
# 3. Get the number of survivors by embarked location
print(data[data['Survived']==1]['Embarked_S'].value_counts())
# 4. Get the number of non-survivors by embarked location
print(data[data['Survived']==0]['Embarked_S'].value_counts())
# 5. Calculate the percentage of children (Age < 18) who
survived
print(data[data['Age']<18]['Survived'].mean())
# 6. Calculate the percentage of adults (Age >= 18) who
survived
print(data[data['Age']>=18]['Survived'].mean())
# 7. Get the median age of survivors
print(data[data['Survived']==1]['Age'].median())
# 8. Get the median age of non-survivors
print(data[data['Survived']==0]['Age'].median())
# 9. Get the median fare of survivors
print(data[data['Survived']==1]['Fare'].median())
# 10. Get the median fare of non-survivors
print(data[data['Survived']==0]['Fare'].median())
```

OUTPUT

Average time
0.035 s
35.00 ms

Maximum time
0.035 s
35.00 ms

1 out of 1 shown test case(s) passed

Test case 1 35 ms Debug

Expected output

female ····233
male ·····109
Name: Sex, dtype: int64
male ·····468
female ·····81
Name: Sex, dtype: int64
1 ····217
0 ····125
Name: Embarked_S, dtype: int64
1 ····427
0 ····122
Name: Embarked_S, dtype: int64
0.5398230088495575
0.3810316139767055
28.0
28.0
26.0
10.5
Actual output

female ····233
male ·····109
Name: Sex, dtype: int64
male ·····468
female ·····81
Name: Sex, dtype: int64
1 ····217
0 ····125
Name: Embarked_S, dtype: int64
1 ····427
0 ····122
Name: Embarked_S, dtype: int64
0.5398230088495575
0.3810316139767055
28.0
28.0
26.0
10.5